

A deep dive into the convolution operation

Romain Thelineau

qwertee.io

romain.thelineau@qwertee.io



@romaintha

PyData Cluj-Napoca, November 25, 2019

About


Ph.D. in physics (Grenoble, France)
Quantum computing hardware



Machine learning engineer
Computer vision



Convolutions are everywhere!



Go to notebook

Convolution: mathematical formalism

The mathematical definition of the convolution of the functions f and g wrtitten $f * g$ is:

$$s(t) = (f * g)(t) = \int_{-\infty}^{-\infty} f(\tau)g(t - \tau)d\tau$$


In the discrete case, the convolution is expressed as:

$$s(t) = (f * g)(t) = \sum_{n=-\infty}^{-\infty} f(n)g(t - n)$$

From 1d to 2d convolutions

Considering an image I and a kernel(filter) K , both being 2 dimensional, the convolution S of I and K is:

$$S(i,j) = (I * K)(i,j) = \sum_{m=-\infty}^{-\infty} \sum_{n=-\infty}^{-\infty} I(m,n)K(i-m,j-n)$$



Convolutions for image processing

Go to notebook

Image recognition



Learn the conditional probability $p(y|x)$ where:
 y : muffin or chihuahua
 x : input image

Deep learning: Universal Approximation Theorem

Universal Approximation Theorem¹:

a MLP can approximate any function with any degree of accuracy (given enough hidden units)






Image source: <http://www.mdpi.com/>

But:


No guarantee of learning it

Number of hidden units required
 $\propto e^{\text{input_size}} \rightarrow \text{Overfitting}$



¹Hornik et al (1989)

Deep learning: MLP




Any output y_i is simply a linear combination of all inputs x_j to which an activation function σ is applied:


$$y_i = \sigma(b_i + \sum_{j=1}^N w_{j,i} x_j)$$

Image key properties


Locality



Translation invariance



Deep learning: convolution



The output y_i can be written as a linear combination of only a few input x_j :

$$y_i = \sigma(b_i + w_1 x_i + w_2 x_{i-1} + w_3 x_{i+1})$$

This leverages 3 important ideas: **sparse connectivity**, **parameter sharing**, **equivariance to translation**

Deep learning: convolution arithmetic

$$o = \left\lceil \frac{i + 2p - k}{s} \right\rceil + 1$$

where:

- o: output size
- i: input size
- p: zero padding (number of zeros added around the image)
- k: kernel size
- s: stride (distance between two consecutive positions of the kernel)

More on convolution arithmetic [here](#).

Deep learning: deep convolutional network

Stacking simple blocks allows for learning higher level features ¹





Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

¹ Lecun et al

Deep learning: visualizing CNN

Visualizing and Understanding Convolutional Networks

Matthew D. Zeiler


Dept. of Computer Science, Courant Institute, New York University

ZEILER@CS.NYU.EDU

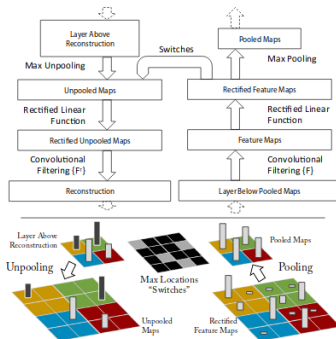
Rob Fergus

Dept. of Computer Science, Courant Institute, New York University

FERGUS@CS.NYU.EDU




Go to notebook



Deep learning: other kind of convolutions

Transposed convolution
(a.k.a. deconvolution)



Dilated convolution
(a.k.a. atrous convolution)




Image by Vincent Dumoulin

- Deep learning book
- Dive into deep learning book
- Understanding convolutions
- Feature visualization
- Deconvolution and checkerboard artifacts

Thank you for your attention